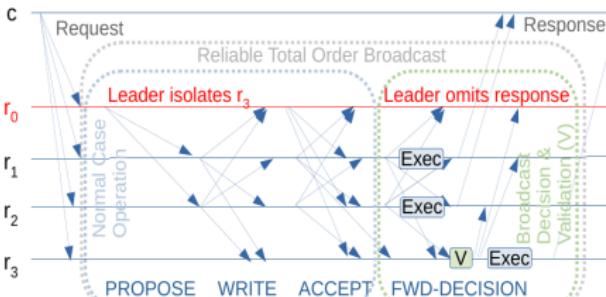


# Dependable Distributed Systems – 5880V/UE

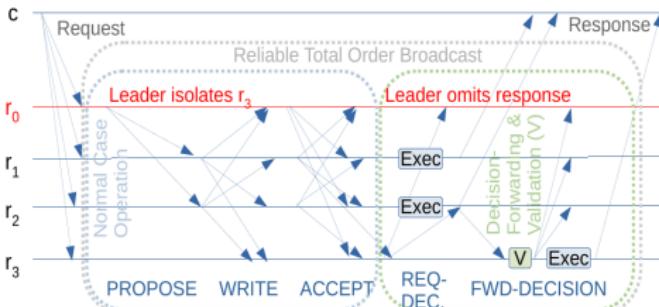
Part 1a: Organization, Introduction – 2021-10-04

Prof. Dr. Hans P. Reiser | WS 2021/22

UNIVERSITÄT PASSAU



This sequence diagram illustrates a normal case operation in a distributed system with four nodes ( $r_0, r_1, r_2, r_3$ ). The timeline is divided into two phases: 'Normal Case Operation' and 'FWD-DECISION'. In the first phase, a 'Request' is broadcast from node  $r_0$ . A leader isolates node  $r_3$ , which then performs a 'PROPOSE' operation. Subsequent 'WRITE' and 'ACCEPT' operations are shown. In the second phase, a 'Reliable Total Order Broadcast' occurs, followed by a 'Leader omits response' and a 'Broadcast & Decision & Validation (V)'. Finally, a 'Response' is sent back to the requester.



This sequence diagram shows a leader omission during a reliable total order broadcast. It follows a similar timeline to the first diagram but includes an additional 'REQ-DEC' step. After the 'FWD-DECISION' phase, the system enters a 'Decision-Forwarding & Validation (V)' phase. During this phase, the leader omits a response, and a broadcast for validation is sent. The process concludes with a 'Response' to the requester.

Faculty of Computer Science and Mathematics – Chair for Reliable Distributed Systems

[www.fim.uni-passau.de/zvs/](http://www.fim.uni-passau.de/zvs/)

# Basic information

- This year: 2-week block lecture, Oct 4 - Oct 14
- Lecture 9-12h, exercises 13-16h  
(but we will flexibly adjust this if needed)
- Master PO2016: DDS → “IT security and reliability”
- Prerequisite:
  - Good basic understanding of computers, networks, operating systems and distributed systems
- Teaching language: English (unless *all* participants agree differently)

## Basic information: Exam

- Oral exam (can be done online)
- First exam: planned for 03.03. and 04.03. (register before 12.12.2021)
- Second exam: planned for 04.04. (register before 25.03.2022)
- Registration for exams on HISQIS
- We will assign time slots in early February
  - Please pay attention to communication on Stud.IP. We may ask you for your preferences for a specific day, online vs. on campus exam, etc.

## Basic information: Lecture material

- Lecture material (slides as PDF document) available on Stud.IP
- Exercise questions will be published on Stud.IP
  - In the exercise classes, YOU shall try to answer all questions on your own
  - ... and then present your solutions

# Basic information

- about:me
  - Hans P. Reiser
  - Office: ITZ 137
  - hr@sec.uni-passau.de
- Secretary: Siglinde Böck, ITZ 130
  - sis-office@sec.uni-passau.de
- Use Stud.IP forum for questions and discussions of potentially public interest!

# Goals of this lecture

In-depth understanding of distributed systems which . . .

- . . . consist of many participants
- . . . have a high level of complexity
- . . . process large amounts of data
- . . . must be fault-tolerant
- . . . must be secure

# Goals of this lecture

In-depth understanding of distributed systems which . . .

- . . . consist of many participants
- . . . have a high level of complexity
- . . . process large amounts of data
- . . . must be fault-tolerant
- . . . must be secure

Main focus: **distributed algorithms**

- A very interesting subject ☺

# Literature

Main material:

- Material available on Stud.IP and ILIAS
- Scientific articles on specific topics

Recommended books:

- C. Cachin, R. Guerraoui, L. Rodrigues:  
Introduction to Reliable and Secure Distributed Programming, Springer, 2011
- M. Raynal:  
Concurrent Programming: Algorithms, Principles, and Foundations, 2013
- P. Veríssimo, L. Rodrigues:  
Distributed Systems for System Architects, Kluwer Academic Publishers, 2001. Part I  
(Distributed systems) and Part II (Fault tolerance)
- Vijay K. Garg: Elements of Distributed Computing, Wiley, 2002.
- Israel Koren, C. Mani Krishna: Fault-Tolerant Systems, Morgan Kaufmann, 2007.

# Semester overview (preliminary plan)

- 1 Basic concepts of dependable distributed systems
- 2 Models, logical and physical clocks and time synchronization
- 3 Network Time Protocol and Precision Time Protocol
- 4 Fault tolerant group communication
- 5 Consensus and failure detectors
- 6 Replicated systems
- 7 Paxos algorithm
- 8 Byzantine fault-tolerant replication
- 9 Blockchain / Distributed Ledger Technology
- 10 Distributed coordination services (Zookeeper) and data storage (GFS etc.)
- 11 Decentralized self-organizing systems (P2P systems)

# Additional remarks regarding the lecture

Feedback and questions:

- Ask question any time (including during lecture!)
- Point out errors!
- Give feedback on the material!
- Feel free to contact me:
  - personally (after lecture, or with appointment in my office)
  - E-mail ([hr@sec.uni-passau.de](mailto:hr@sec.uni-passau.de))

# Basic information

Questions regarding organization?

# Overview

## 1 Organisation

## 2 Introduction

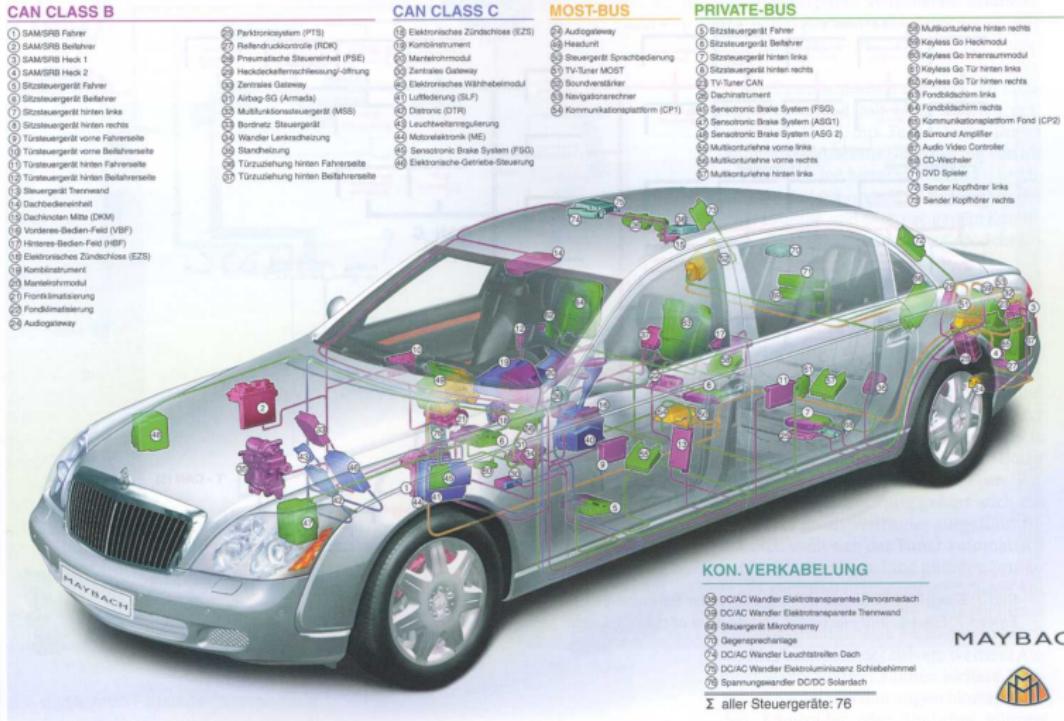
- Distributed systems
- Basics of fault tolerance
- Ways towards dependability

# Distributed Systems: Fundamentals

## Overview

- What is a distributed system?
- Properties of a distributed system

# What is a distributed system? (Example)



(source: Der neue Maybach, ATZ/MTZ Sonderheft September 2002, Seite 125)

# What is a distributed system? (Example)



(source: NapoliRoma, CC BY-SA 3.0, [https://commons.wikimedia.org/wiki/File:Sun\\_Modular\\_Datacenter.jpg](https://commons.wikimedia.org/wiki/File:Sun_Modular_Datacenter.jpg))

Organisation

Distributed systems

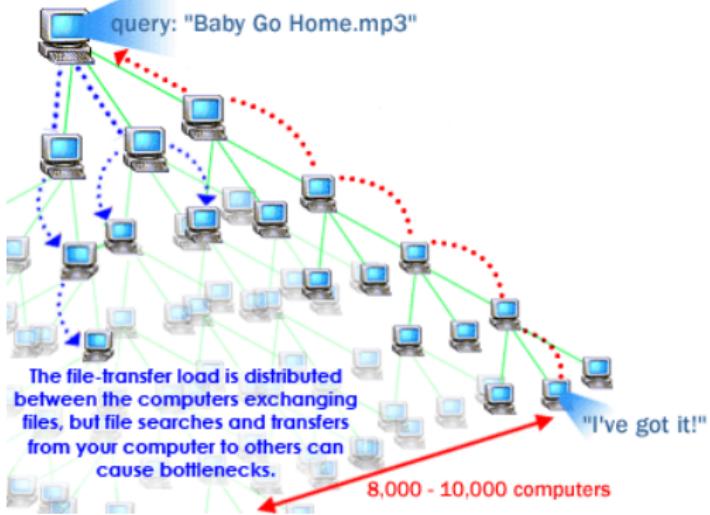
Basics of fault tolerance

Introduction

Ways towards dependability

# What is a distributed system? (Example)

© 2002 HowStuffWorks



(source: Carmen Carmack "How BitTorrent Works" 26 March 2005. <https://computer.howstuffworks.com/bittorrent.htm>)

(source: Sandvine, 2012 (North America, Fixed Access))

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	36.8%	Netflix	33.0%	Netflix	28.8%
2	HTTP	9.83%	YouTube	14.8%	YouTube	13.1%
3	Skype	4.76%	HTTP	12.0%	HTTP	11.7%
4	Netflix	4.51%	BitTorrent	5.89%	BitTorrent	10.3%
5	SSL	3.73%	iTunes	3.92%	iTunes	3.43%
6	YouTube	2.70%	MPEG	2.22%	SSL	2.23%
7	PPStream	1.65%	Flash Video	2.21%	MPEG	2.05%
8	Facebook	1.62%	SSL	1.97%	Flash Video	2.01%
9	Apple PhotoStream	1.46%	Amazon Video	1.75%	Facebook	1.50%

Organisation

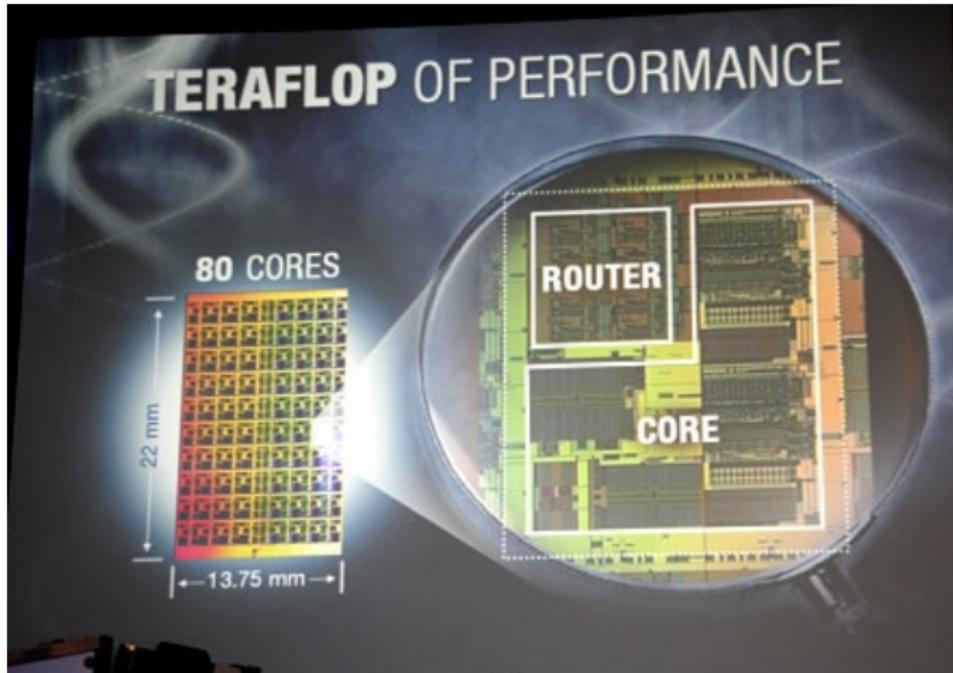
Distributed systems

Basics of fault tolerance

Introduction

Ways towards dependability

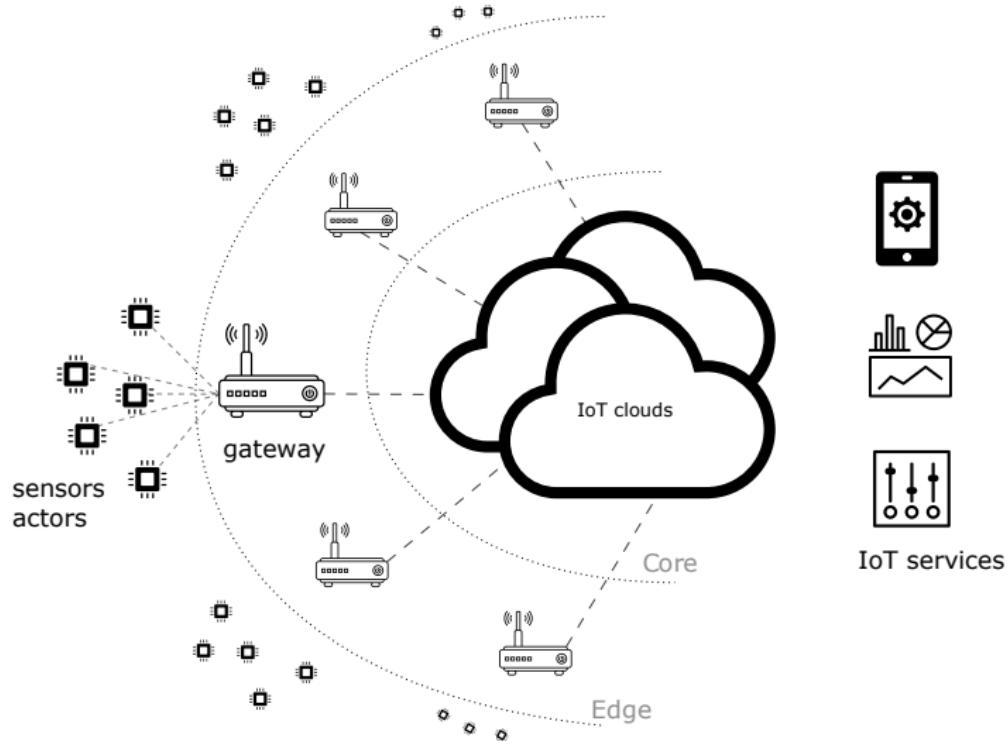
# What is a distributed system? (Example)



(Completely different dimension, but problems similar to the other examples)

(source: Intel, Intel Developer Forum Fall 2006)

# What is a distributed system? (Example)



# What is a distributed system?

- A computer network is **not** distributed system

# What is a distributed system?

- A computer network is **not** distributed system
- Computer network
  - Infrastructure that enables communication between machines
  - May have multiple media and topologies
  - Common communication protocols
- Distributed system
  - Composed of multiple computer systems
  - Communicate over a computer network
  - Set of processes that
    - use a common set of distributed protocols
    - to support a coherent execution
    - of distributed activities

# Properties of a distributed system

- Several computers
- Connected by communication network ⇒
- Distributed global state

# Properties of a distributed system

- Independent failures
- Several computers
- Connected by communication network ⇒
- Distributed global state

# Properties of a distributed system

- Several computers
  - Connected by communication network
  - Distributed global state
- ⇒
- Independent failures
  - Unreliable communication

# Properties of a distributed system

- Several computers
    - Independent failures
    - Unreliable communication
    - Variable delays
  - Connected by communication network
  - Distributed global state
- ⇒

# Properties of a distributed system

- Several computers
  - Independent failures
  - Unreliable communication
  - Variable delays
  - Moderate communication speed
- Connected by communication network
- ⇒
- Distributed global state

# Properties of a distributed system

- Several computers
  - Connected by communication network
  - Distributed global state
- ⇒
- Independent failures
  - Unreliable communication
  - Variable delays
  - Moderate communication speed
  - Only partial ordering of events

# Properties of a distributed system

- Several computers
  - Connected by communication network
  - Distributed global state
- ⇒
- Independent failures
  - Unreliable communication
  - Variable delays
  - Moderate communication speed
  - Only partial ordering of events
  - Global state difficult to determine

# Properties of a distributed system

- Several computers
- Connected by communication network
- Distributed global state
- Independent failures
- Unreliable communication
- Variable delays
- Moderate communication speed
- Only partial ordering of events
- ⇒ ■ Global state difficult to determine
- Higher administrative costs

# Properties of a distributed system

- Several computers
- Connected by communication network
- Distributed global state
- Independent failures
- Unreliable communication
- Variable delays
- Moderate communication speed
- Only partial ordering of events
- ⇒ ■ Global state difficult to determine
- Higher administrative costs
- Is often less expensive than centralized supercomputers

# Distributed system: definitions

- “A distributed system is a collection of independent computers that appear to the users of the system as a single computer.”

*Andrew Tanenbaum*

# Distributed system: definitions

- No common clock
  - Distributed local clocks can be synchronized with only limited accuracy
- No shared memory
  - Common state not stored centrally;  
observing global system properties difficult
- No exact failure detection
  - No upper bounds on communication delays;  
failure not distinguishable from slow node

Vijay Garg

# Centralized vs. distributed systems

## Centralized:

- Easy accessibility
  - of resources and information
- Homogeneity
  - of methods and technologies
- Manageability
- Consistency
  - a global state
- Security
  - easier to protect

## Distributed:

- Scalability, by
  - size
  - modularity
- Reliability and availability
  - redundancy and replication
  - graceful degradation
- Cheaper
- Security
  - by being able to compensate the impact of security problems
  - not really easy

# Large Data Center (I)

Google data center on the banks of the Columbia River. The site, with 2 server-packed buildings and space for a 3rd, houses tens of thousands of computers. Microsoft, Yahoo, and Amazon are also building data centers in the region.

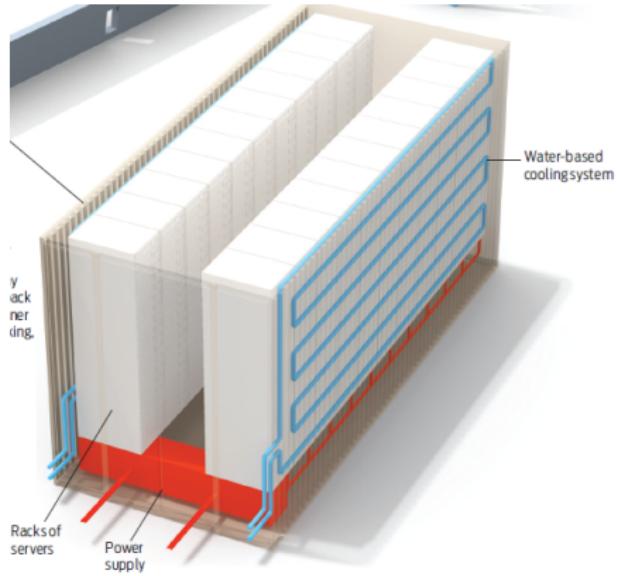
- Randy Katz, Tech Titans Building Boom, IEEE Spectrum, Feb. 09



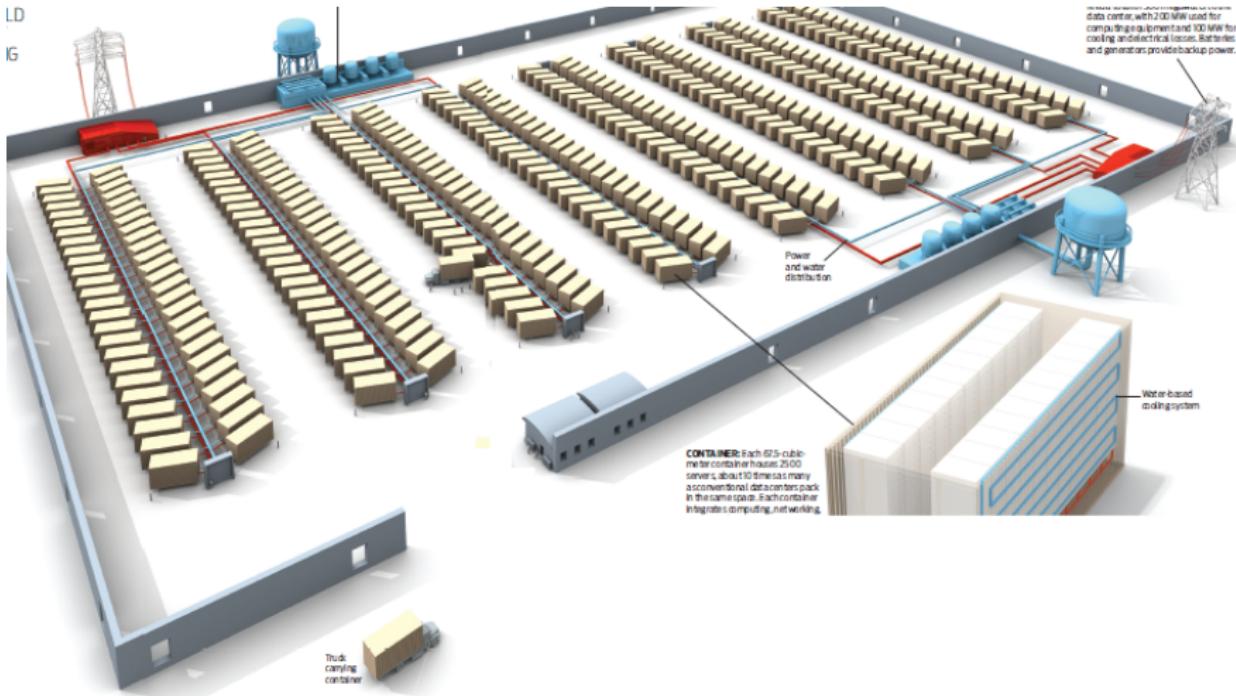
# Large Data Center (II)

## CONTAINER:

- Each 67.5-cubic-meter container houses 2500 servers, about 10 times as many as conventional data centers pack in the same space.
- Each container integrates computing, networking, power, and cooling systems.



# Large Data Center (III)



# Large Data Center (IV)

## The Joys of Real Hardware

Typical first year for a new cluster:

- ~0.5 **overheating** (power down most machines in <5 mins, ~1-2 days to recover)
- ~1 **PDU failure** (~500-1000 machines suddenly disappear, ~6 hours to come back)
- ~1 **rack-move** (plenty of warning, ~500-1000 machines powered down, ~6 hours)
- ~1 **network rewiring** (rolling ~5% of machines down over 2-day span)
- ~20 **rack failures** (40-80 machines instantly disappear, 1-6 hours to get back)
- ~5 **racks go wonky** (40-80 machines see 50% packetloss)
- ~8 **network maintenances** (4 might cause ~30-minute random connectivity losses)
- ~12 **router reloads** (takes out DNS and external vips for a couple minutes)
- ~3 **router failures** (have to immediately pull traffic for an hour)
- ~dozens of minor **30-second blips for dns**
- ~1000 **individual machine failures**
- ~thousands of **hard drive failures**
- slow disks, bad memory, misconfigured machines, flaky machines, etc.**

Long distance links: **wild dogs, sharks, dead horses, drunken hunters, etc.**

source: Jeffrey Dean, Google, <https://research.google/people/jeff/>

# Overview

## 1 Organisation

## 2 Introduction

- Distributed systems
- **Basics of fault tolerance**
- Ways towards dependability

# Fault Tolerance: Basics

Why fault tolerance?

# Computer failures

*Why do computers stop and what can be done about it? (J. Gray, 1985)*

# Computer failures

*Why do computers stop and what can be done about it? (J. Gray, 1985)*

- Because:
  - Everything that works will fail at some time
  - Hardware and software is generally overestimated
- Better prevent errors instead of repairing later
  - In a predictable and repeatable manner
- We need:
  - a scientific way to *quantify, predict, prevent and tolerate* disturbances that disrupt the normal operation of a system

# Definition

Laprie 1993

“Computer system dependability is the quality of a delivered service such that reliance can justifiably be placed on this service”

# Definition

Laprie 1993

“Computer system dependability is the quality of a delivered service such that reliance can justifiably be placed on this service”

A reliable system is a system that with high probability works according to its specification

- A simple definition . . .

# Definition

Laprie 1993

“Computer system dependability is the quality of a delivered service such that reliance can justifiably be placed on this service”

A reliable system is a system that with high probability works according to its specification

- A simple definition . . .  
. . . but some open questions:
  - How is the service specified?
  - What exactly is specified?  
(the service itself? the environment?)
  - What is a “high probability”?

# Impact of distribution on dependability?

“A distributed system is one in which the failure of a computer you did not even know existed can render your own computer unusable.”

*Leslie Lamport*

# Impact of distribution on dependability?

"A distributed system is one in which the failure of a computer you did not even know existed can render your own computer unusable."

*Leslie Lamport*

- A failure of some component might influence other components!
- Communication usually happens
  - over unreliable networks
  - with unpredictable delays

# Impact of distribution on dependability?

- Failure of a group of computers:
  - If a single component fails on average once per year...  
(and all components fail independently)
  - ... how many failures per year will you have in a datacenter using 1000 of these components?

## Malicious faults (attackers!) make it even worse

- Errors are no longer independent
- Errors are more serious
- Error models are inaccurate  
(... stochastic model of an attacker?)

# Dependability attributes

Basic definitions: (source: DSSA)

- Reliability
  
- Maintainability
  
- Availability

# Dependability attributes

Basic definitions: (source: DSSA)

- Reliability
  - Measure of the continuous delivery of correct service
  - Metrics:
    - “Mean Time To Failure” (MTTF)
    - “Mean Time Between Failures” (MTBF)
    - Probability that the system does not fail during mission period
    - Probabilistic failure rate (e.g.,  $10^{-9}$  failures/hour)
- Maintainability
- Availability

# Dependability attributes

Basic definitions: (source: DSSA)

- Reliability
  - Measure of the continuous delivery of correct service
  - Metrics:
    - “Mean Time To Failure” (MTTF)
    - “Mean Time Between Failures” (MTBF)
    - Probability that the system does not fail during mission period
    - Probabilistic failure rate (e.g.,  $10^{-9}$  failures/hour)
- Maintainability
  - Measure of the time to restoration of correct service
  - Metric: “Mean Time To Repair (MTTR)”
- Availability

# Dependability attributes

Basic definitions: (source: DSSA)

- Reliability

- Measure of the continuous delivery of correct service
- Metrics:
  - “Mean Time To Failure” (MTTF)
  - “Mean Time Between Failures” (MTBF)
  - Probability that the system does not fail during mission period
  - Probabilistic failure rate (e.g.,  $10^{-9}$  failures/hour)

- Maintainability

- Measure of the time to restoration of correct service
- Metric: “Mean Time To Repair (MTTR)”

- Availability

- Measure of the delivery of correct service  
with respect to the alternation  
between correct and incorrect service
- Availability =  $MTBF / (MTBF+MTTR)$

# Other properties

- Integrity
  - Absence of improper system alterations
- Security
  - Ability to guarantee confidentiality, integrity and availability in service provision
- Resilience
  - Ability to recover from external faults
  - Frequently used when considering applications that require interactions through their surrounding environment

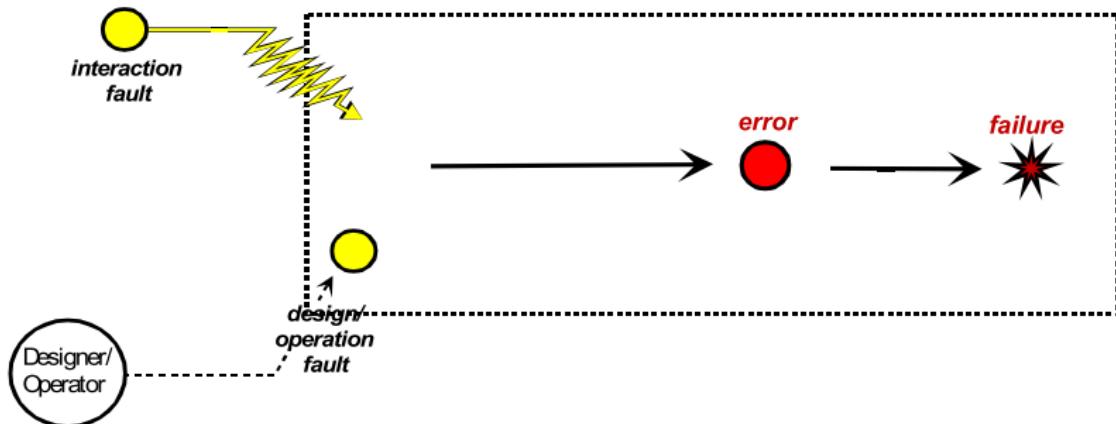
# Faults, errors and failures

- A system **failure** occurs when the delivered service deviates from fulfilling the system function
- An **error** is that part of the system state which is liable to lead to subsequent failure
- The adjudged cause of an error is a **fault**
- These notions are for a system!

# Faults, errors and failures

- A system **failure** occurs when the delivered service deviates from fulfilling the system function
- An **error** is that part of the system state which is liable to lead to subsequent failure
- The adjudged cause of an error is a **fault**
- These notions are for a system!
  
- Example (system = RAM memory)
  - Fault — stuck-at-'0' RAM memory register
  - Error — state after '1' is written to RAM
  - Failure — read returns wrong value '0'

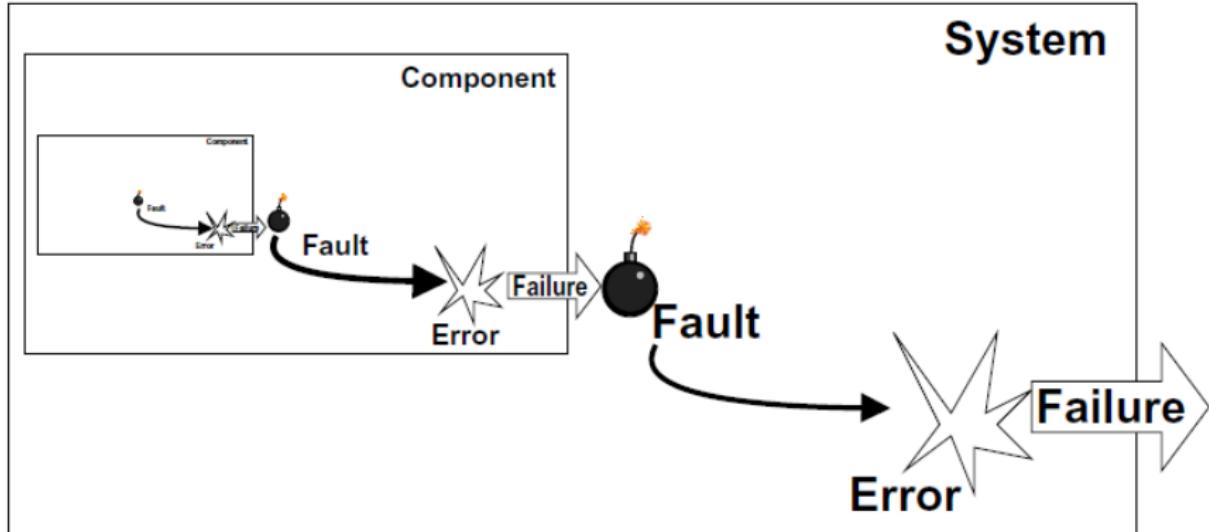
# Sequence fault – error – failure



(source: Paulo Veríssimo)

- The relation is recursive!
  - ... What is a system? What is a component?

# Fault – error – failure recursiveness



(source: Paulo Veríssimo)

# Today's lecture

## 1 Organisation

## 2 Introduction

- Distributed systems
- Basics of fault tolerance
- Ways towards dependability

# Ways towards dependability

- **Fault prevention** — how to prevent the occurrence of faults?
  - Avoid design defects (good process; tools; etc.)
  - Shielding, robust design, good operational procedures etc. for runtime faults
- **Fault tolerance** — how to deliver correct service in the presence of faults?
- **Fault removal** — how to reduce the number/severity of faults?
  - Verification & validation (design phase)
  - Corrective & preventive maintenance (operations)
- **Fault forecasting** — how to estimate the future incidence of faults?
  - Simulation, modeling, prediction (data from models)
  - Historical data, accelerated life testing, etc. (data from real systems)

# Ways towards dependability

- **Fault prevention** — how to prevent the occurrence of faults?
  - Avoid design defects (good process; tools; etc.)
  - Shielding, robust design, good operational procedures etc. for runtime faults
- **Fault tolerance** — how to deliver correct service in the presence of faults?
  - $\Rightarrow$  Main topic of lecture
- **Fault removal** — how to reduce the number/severity of faults?
  - Verification & validation (design phase)
  - Corrective & preventive maintenance (operations)
- **Fault forecasting** — how to estimate the future incidence of faults?
  - Simulation, modeling, prediction (data from models)
  - Historical data, accelerated life testing, etc. (data from real systems)

# Fault tolerance

Elements of fault tolerance:

- Error detection
- Recovery
  - Error handling
    - Backward recovery
    - Forward recovery
    - Compensation/masking
  - Fault handling

# Error detection

- Reactive detection

- Detects and reports errors at run-time (during operation)
- Parity bits on memory, data buses, ALU, etc.
- Coverage and flow-checking
  - Did the software hit all required checkpoints?
  - Multi-version comparison and majority voting
  - Compute multiple times and compare

# Error detection

- Reactive detection
  - Detects and reports errors at run-time (during operation)
  - Parity bits on memory, data buses, ALU, etc.
  - Coverage and flow-checking
    - Did the software hit all required checkpoints?
    - Multi-version comparison and majority voting
    - Compute multiple times and compare
- Proactive detection
  - Detects and reports errors at development-time (offline)
  - Idea is to activate any faults before they can affect real computations
  - Execute self-tests, data consistency checks
  - Exercise system to activate faults
    - Memory scrubbing daemon that reads/resets all memory locations periodically

# Error/fault recovery

- The objective is to restore a system state with errors and (possibly) faults into one without errors and faults
- Recovery can be done by:
  - Fault handling
  - Error handling
    - Backward recovery
    - Forward recovery
    - Error compensation or error masking

# Fault handling

Record faults to avoid later reactivation

- Fault diagnosis: identify cause
  - Once diagnosed, avoid using that resource/performing that operation again
  - Also known as the “Don’t do that” solution
- Fault isolation: eliminate access to source of fault
  - Exclude faulty component from use
  - Shut down defective server in server farm
- System reconfiguration
  - Switch in spares or reassign tasks
  - Route traffic to backup component
- System reinitialization
  - Rebooting after repair/reconfiguration
  - Rebooting to perform “software rejuvenation”

# Error handling

- Roll-back: return system to its state before the error occurred
  - Database commit (all-or-nothing execution semantics)
    - If error, undo all partial results to restore clean system state
  - Checkpoint/rollback
    - Periodically snapshot system state and restore upon error to avoid failure
    - Reboot is a very drastic rollback (reboot) to a checkpoint (freshly system image)
- Roll-forward: move the system forward into an error-free state
  - Retry to obtain error-free result
  - Proceed on and let error flush itself out of system
    - Replication where a backup takes over after the working replica fails

# Error masking

- How?

- The system state has enough **redundancy** that the correct service can be provided without any noticeable glitch
- Recovery is achieved without explicit error detection

- What's good?

- User is blissfully unaware of anything bad happening
- Fault/failure transparency — great way to build systems!

- What's bad?

- Independent failure assumption
- Does not deal adequately with design faults

# Forms of redundancy

- Space redundancy
  - Several copies of the same component
  - Same information stored in several disks
  - Different nodes compute same result in parallel
  - Messages disseminated along different paths
- Time redundancy
  - Doing the same thing more than once, in same or different ways
  - Retransmission of lost messages
  - Repeating computations that have aborted
- Value redundancy
  - Adding extra information about the data being stored or sent
  - Codes that allow the detection or correction of integrity errors
  - Parity bit or ECC added to memory chips or disk structures
  - Frame check sequences or cyclic redundancy in transmitted data
  - cryptographic message signatures

# The “Dependability Tree”

## IMPAIRMENTS

- FAULTS
- ERRORS
- FAILURES

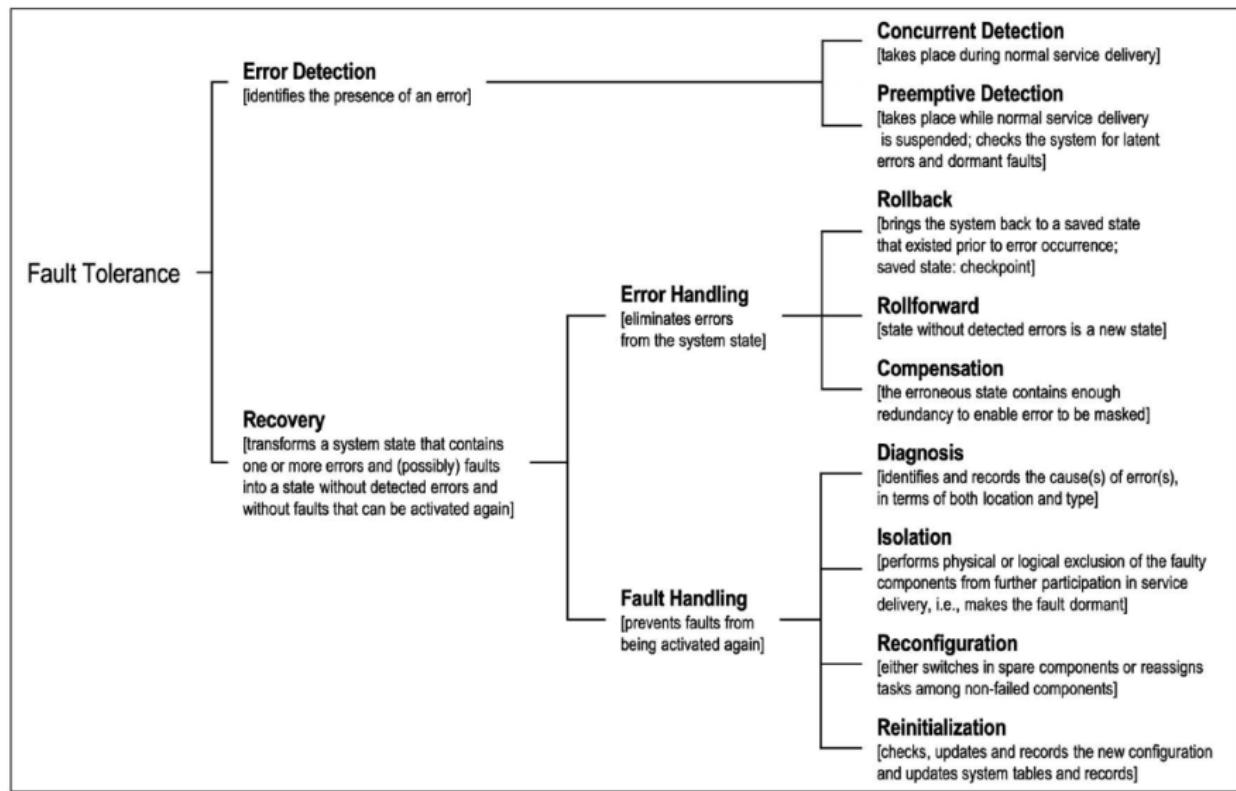
## MEANS

- FAULT PREVENTION
- FAULT TOLERANCE
- FAULT REMOVAL
- FAULT FORECASTING

## ATTRIBUTES

- RELIABILITY
- MAINTAINABILITY
- AVAILABILITY
- SAFETY
- CONFIDENTIALITY
- INTEGRITY
- AUTHENTICITY

# The “Fault Tolerance Tree”



Avizienis et al.: Basic concepts and taxonomy of dependable and secure computing, TDSC v1(1), March 2004

# Resilience vs. dependability

## Resilience definition (Laprie)

The persistence of service delivery that can justifiably be trusted, when facing changes.

J.-C. Laprie: "From dependability to resilience", DSN, 2008

## Resilience definition (Berger et al.)

Resilience is the property of preserving the dependability and security of a system when the system encounters changes, thus withstanding or recovering from impairments.

Berger et al., "A Survey on Resilience in the IoT: Taxonomy, Classification and Discussion of Resilience Mechanisms", ACM CSUR, 2021

# Summary

Basic concepts you should remember

- What characterizes a distributed system?
  - Properties and problems
- What is dependability?
  - Dependability attributes
  - Faults, errors, and failures
  - Ways towards dependability