


Dependable Distributed Systems – 5880V/UE

Part 1b: Distributed Systems: Models, Time – 2021-10-04

Prof. Dr. Hans P. Reiser | WS 2021/22

UNIVERSITÄT PASSAU



This week's lecture

- 1 Models of distributed systems
 - Modelling processes
 - Modelling communication
 - Modelling time
- 2 Relevance of time in distributed systems
 - Time and the order of events
 - Problem of clock synchronization
- 3 Logical clocks
 - Lamport clocks
 - Vector clocks

This week's lecture

1 Models of distributed systems

- Modelling processes
- Modelling communication
- Modelling time

2 Relevance of time in distributed systems

- Time and the order of events
- Problem of clock synchronization

3 Logical clocks




- Lamport clocks
- Vector clocks

Models

- Basic idea: a *model* is a simplification of an object (= system)
 - Allows us to reason about it
- A *model* for an object is a collection of attributes and a set of rules that govern how these attributes interact
 - Also called a theory
- Two important facts:
 - There is no single correct model
 - Answering different types of questions usually requires different models

Models of the atom

- Billiard Ball Model (1803) – John Dalton: atom: small solid sphere
- Plumb Pudding Model (1897) – Joseph John Thomson
- Solar System Model – Ernest Rutherford (1911), Neils Bohr (1913)
- Electron Cloud Model (1920's)



Model \Rightarrow Assumptions

- Two processes, A and B, communicate by sending and receiving messages on a bidirectional channel. Neither process can fail. However, the channel can experience transient failures, resulting in the loss of a subset of the messages that have been sent.

Model \Rightarrow Assumptions

- Two processes, A and B, communicate by sending and receiving messages on a bidirectional channel. Neither process can fail. However, the channel can experience transient failures, resulting in the loss of a subset of the messages that have been sent.

What are the **assumptions** above?

Good models

- A model is *accurate* to the extent that analyzing it yields truths about the object of interest
- A model is *tractable* if such an analysis is actually possible

Good models

- A model is *accurate* to the extent that analyzing it yields truths about the object of interest
- A model is *tractable* if such an analysis is actually possible
- Defining an accurate model is not difficult; defining an accurate+tractable model is
 - An accurate+tractable model will include exactly those attributes that affect the phenomena of interest
 - Level of detail is a key issue

Good models

- In building models for distributed systems, we typically seek answers to two fundamental questions:
- **Feasibility.** What classes of problems can be solved?
 - Can head-off wasted effort in design, implementation, and testing
- **Cost.** For those classes that can be solved, how expensive must the solution be?
 - Support evaluating any solution we devise
 - Avoid designs requiring protocols that are inherently slow or expensive

Distributed systems

- Distributed systems: models about processes
- Multiple processes communicating over narrow bandwidth, high-latency channels, with some processes and channels faulty

This week's lecture

1 Models of distributed systems

- Modelling processes
- Modelling communication
- Modelling time

2 Relevance of time in distributed systems

- Time and the order of events
- Problem of clock synchronization


3 Logical clocks

- Lamport clocks
- Vector clocks

Modelling processes

- Node, computer, process (A, B, C, \dots or P_1, P_2, P_3, \dots):
 - Used as synonyms
 - Independent, active instances in the distributed systems
 - Interaction with other participants with messages (usually)
- Global state
 - Vector $S = [S_1, S_2, \dots, S_n]$ of local states S_i of all nodes
 - (sometimes: state of communication channels also relevant)
 - We can easily use S in formal arguments, but it is distributed over all nodes, and as such it is not easily observable
- *Step*: atomic local transition of a process P_i from state S_i to S'_i
 - Designated with s, σ (step) or e (event)
 - Spontaneous or triggered by message

Time-space diagram



Modelling faulty processes

- There are many failure models for distributed systems; all based on assigning responsibility for faulty behaviour to system components:
 - processors
 - communications channels

Modelling faulty processes

- There are many failure models for distributed systems; all based on assigning responsibility for faulty behaviour to system components:
 - processors
 - communications channels
- We count *faulty components*, not occurrences of faulty behavior
 - In classical work on fault-tolerant computing systems: occurrences of faulty behaviour are counted
- t -fault tolerant
 - system satisfies its **specification**, provided that **no more than t** of its components are **faulty**

(we also use f or k instead of t , and n for the total number of nodes)

Faulty processes

■ Crash

- A process fails by *stopping* (does not execute any further steps). The process will remain in the crash state forever.

■ Fail-stop

- All other operational processes reliably detect the crash

■ Fail-silent

- Other processes are not reliably notified about the crash
- If the communication system is not synchronous (see later), it is hard to distinguish a crashed from a slow process

Faulty processes

■ Crash

- A process fails by *stopping* (does not execute any further steps). The process will remain in the crash state forever.
- **Fail-stop**
 - All other operational processes reliably detect the crash
- **Fail-silent**
 - Other processes are not reliably notified about the crash
 - If the communication system is not synchronous (see later), it is hard to distinguish a crashed from a slow process

■ Crash-Recovery

- A correct process can fail and recover again (for a finite number of times)
- A process is faulty if it stops forever, or if it infinitely often fails and recovers.
- Special case “**omission**” (fault causes only loss of messages, not loss of local state)

Faulty processes

■ Eavesdropping faults

- Faulty processes behave as in the crash-recovery model (i.e., interactions with other processes occur according to specification or not at all)
- In addition: Faulty processes may pass internal information to an external third party

Faulty processes

■ Eavesdropping faults

- Faulty processes behave as in the crash-recovery model (i.e., interactions with other processes occur according to specification or not at all)
- In addition: Faulty processes may pass internal information to an external third party


■ Byzantine faults

- Faulty processes can exhibit arbitrary behaviour
- Usually: “*arbitrary*” restricted to *computationally feasible operations*
 - Example: cannot break strong cryptography

Important: these are models \Rightarrow assumptions!

Hierarchy of fault classes


Hierarchical relations (according to RSDP):



Severity of faults

Classification according to DSSA:

- More detailed differentiation of some faults:
- *Assertive faults* / Value faults
 - Data in protocol-conforming interactions modified syntactically or semantically
- *Omissive faults*
 - Actions do not take place, or too late/too early



Overview

1 Models of distributed systems

- Modelling processes
- **Modelling communication**
- Modelling time

2 Relevance of time in distributed systems

- Time and the order of events
- Problem of clock synchronization

3 Logical clocks

- Lamport clocks
- Vector clocks

Modelling point-to-point communication

Point-to-point communication



- Sending message m
 - Process p sends a message m to B by executing `send(m)`
 - `send(m)` adds message m to p 's send queue
- Channel transports m
 - from p 's send queue
 - to q 's receive queue
- Reception of message m
 - q 's runtime system calls `deliver(m)`, which
 - delivers the message m to process q , and
 - removes m from the receive queue.

Communication faults

Modelling message loss

- “Fair-loss links”: weak but useful model
 - No infinite message losses (if periodically repeating message, it eventually will arrive)
 - No infinite number of message duplications
 - No spurious messages (delivery of messages that have not been sent)
- “Stubborn links”:
 - Message is repeated infinitely often over fair-loss link
 - Results in reliable delivery without spurious messages
- “Perfect links” / “reliable links”
 - No loss of messages
 - No duplication
 - No spurious messages
- “Authenticated perfect links”
 - Perfect link with correct sender identification

This week's lecture

1 Models of distributed systems

- Modelling processes
- Modelling communication
- **Modelling time**

2 Relevance of time in distributed systems

- Time and the order of events
- Problem of clock synchronization

3 Logical clocks

- Lamport clocks
- Vector clocks

Synchronous versus asynchronous systems

- *Asynchronous* system: we make no assumptions about process execution speeds and/or message delivery delays
- *Synchronous* system: we do make assumptions about these parameters
 - The relative speeds of processes are assumed to be bounded
 - The delays associated with communications channels also are

Synchronous versus asynchronous systems

- *Asynchronous* system: we make no assumptions about process execution speeds and/or message delivery delays
- *Synchronous* system: we do make assumptions about these parameters
 - The relative speeds of processes are assumed to be bounded
 - The delays associated with communications channels also are
- Postulating that a system is asynchronous is a non-assumption
 - **Every** system “is” asynchronous (i.e. satisfies assumptions that are made)
 - Algorithm for asynchronous system . . . can be used on every system!

Synchronous versus asynchronous systems

- Postulating that a system is synchronous constrains how processes and communications channels are implemented
 - Scheduler that multiplexes processors must not violate the constraints on process execution speeds
 - This implies that all processors in the system have access to approximately rate-synchronized real-time clocks
 - Queuing delays, unpredictable routings, and retransmission due to errors must not violate the constraints on channel delays

Synchronous versus asynchronous systems

- Postulating that a system is synchronous constrains how processes and communications channels are implemented
 - Scheduler that multiplexes processors must not violate the constraints on process execution speeds
 - This implies that all processors in the system have access to approximately rate-synchronized real-time clocks
 - Queuing delays, unpredictable routings, and retransmission due to errors must not violate the constraints on channel delays
 - In asserting that a system is synchronous, we rule out certain system behaviors
 - This enables us to employ
 - simpler protocols (complexity)
 - cheaper protocols (overhead)
- than required in an asynchronous system

Synchronous versus asynchronous systems

There are other models in the spectrum:

- Partial synchrony
- Timed-asynchronous
- Wormholes
- . . .

More details see later lectures. . .

Wrap-up of this part

What you should remember about models and distributed systems:

What you should remember about models:

- How to model processes and process failures
- How to model communication in a distributed system
- How to model time in a distributed system

Time in distributed systems

1 Models of distributed systems

- Modelling processes
- Modelling communication
- Modelling time

2 Relevance of time in distributed systems

- Time and the order of events
- Problem of clock synchronization

3 Logical clocks

- Lamport clocks
- Vector clocks

Time and the order of events: make example

```
% vi test.c  
[...]
```


```
% make  
make: Warning: File `test.c' has modification time 94 s in the future  
cc test.c -o test  
make: warning: Clock skew detected. Your build may be incomplete.
```

```
% make  
make: Warning: File `test.c' has modification time 93 s in the future  
cc test.c -o test  
make: warning: Clock skew detected. Your build may be incomplete.
```

- Make tool does not recognize that *test.c was not modified. . .*
- Or worse: modification is not detected, *file will not get compiled!*

Time and the order of events: Checkpoints

- Given: A distributed system of interacting nodes
- For a consistent checkpoint, the local state of a set of nodes shall be saved at a fixed time.
- Problem: inconsistency may occur if the local clocks differ (C3)



C3: Effect of m is included in S_2 , but not in S_1 !

Time and the order of events

Examples: monitoring, debugging, determining the root causes of errors

- Distributed data acquisition for logging or debugging purposes
- Real order cannot be reconstructed if locally generated timestamps are not synchronized

Problem of clock synchronization

- There are no completely identical physical clocks
 - Different initialization (constant offset)
 - Different speed (frequency error)
 - Subject to environmental conditions (e.g., component ageing, temperature dependency)
- ⇒ Without synchronization errors may continuously grow!
- Common clock for all the nodes of a distributed system (usually) not feasible
- Central reference clocks
 - e.g., radio transmission (WWV, DCF77, GPS)with technically limited accuracy

Time in distributed systems

1 Models of distributed systems

- Modelling processes
- Modelling communication
- Modelling time

2 Relevance of time in distributed systems

- Time and the order of events
- Problem of clock synchronization

3 Logical clocks

- Lamport clocks
- Vector clocks

Logical clocks

The basic idea (Lamport's logical clock):

- Order of events on different components is relevant
 - only if one event could influence another, i.e.,
 - only if the components interact

Logical clocks

The basic idea (Lamport's logical clock):

- Order of events on different components is relevant
 - only if one event could influence another, i.e.,
 - only if the components interact
- Synchronization upon interaction (i.e., communication), such that
 - timestamps respect real order of events

Logical clocks

Description of the algorithm

- Each process P_i has a local clock C_i that counts as follows:

- Local action or send action by P_i : $C_i = C_i + 1$

timestamp of the action: value after increasing

- Message m carries a timestamp t_m
 t_m equals the timestamp of the send action


- P_i with local clock value C_i
receives a message m with timestamp t_m :

$$C_i = \max(C_i, t_m) + 1$$

Resulting C_i : timestamp of the reception event

Logical clocks

Example:



Logical clocks

Properties

- Causal dependency of an event e_2 on event e_1 (i.e., e_1 may have influenced e_2 , shorthand $e_1 \rightarrow e_2$):
$$e_1 \rightarrow e_2 \Rightarrow t(e_1) < t(e_2)$$
- The other direction $t(e_1) < t(e_2) \Rightarrow e_1 \rightarrow e_2$ does not hold!
- The logical timestamps create a partial order on the set of all events!

Logical clocks

Properties

- Causal dependency of an event e_2 on event e_1
(i.e., e_1 may have influenced e_2 , shorthand $e_1 \rightarrow e_2$):
$$e_1 \rightarrow e_2 \Rightarrow t(e_1) < t(e_2)$$
- The other direction $t(e_1) < t(e_2) \Rightarrow e_1 \rightarrow e_2$ does not hold!
- The logical timestamps create a partial order on the set of all events!

Extensions to obtain a total order

- Timestamp of event at process i :
(local time C_i , process number i)
- Order relation:
$$(C_i, i) < (C_k, k) \Leftrightarrow C_i < C_k \vee (C_i = C_k \wedge i < k)$$

Logical clocks: vector clocks

Vector clocks:

- Each process has a local clock, which consists of a vector of N values (N = number of existing nodes)
 - Implementation:
 - Initialization of each vector with the zero vector
 - Between two local events on a process P_i , the corresponding component of the vector is incremented: $C_i[i] = C_i[i] + 1$
 - Messages carry a timestamp t equal to the vector C_i of the sender
 - A process P_i that receives a message
 - incrementes own component in time vector
 - combines it with the received time vector t
- $$C_i[i] := C_i[i] + 1$$
- $$\text{For } k = 1 \dots N : C_i[k] := \max(C_i[k], t[k])$$

Logical clocks: vector clocks

Properties:

- Creates a causal order: $t(e_1) < t(e_2) \Leftrightarrow e_1 \rightarrow e_2$
- Definition of “<”:
 $t(e_1) = (a_1, a_2, \dots, a_n)$ and $t(e_2) := (b_1, b_2, \dots, b_n)$:
 $t(e_1) < t(e_2) \Leftrightarrow (\forall i : a_i \leq b_i) \wedge (\exists i : a_i < b_i)$

Advantage:

- Exact statement about causal relations of events

Disadvantage:

- High communication overhead
 - vector of N elements in each message
 - scales poorly for large N

Summary

Modelling distributed systems

- Processes (including failures), communication, time

Time

- Relevance of time
- Logical clocks: Lamport clocks, vector clocks