


Dependable Distributed Systems – 5880V/UE

Part 3: Broadcast semantics and algorithms – 2021-10-06

Prof. Dr. Hans P. Reiser | WS 2021/22

UNIVERSITÄT PASSAU



Overview

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- Reliable broadcast (rbcast)
- FIFO broadcast (fbcast)
- Causal broadcast (cbcast)
- Atomic broadcast (abcast)

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches

Safety and liveness [according to Lamport]

■ Safety (*Sicherheit*)

■ **Something bad does not happen**

- Safety property defines that some events must *never* happen
- In other words: Safety property is a predicate P that is always true

■ Liveness (*Lebendigkeit*)

■ **Something good will happen**

- Liveness property defines good events that at some time (“eventually”) must happen
- In other words: Liveness property is a predicate P that eventually becomes true

Safety and liveness

Beware of false friends

“eventually” is (close to) the opposite of the german word “eventuell”

- “*X eventually happens*” means that:
- “*X **will happen** within an unspecified but finite amount of time*”

Safety and liveness

Example: Factorization algorithm $Fac(n) \rightarrow \{p_i\}$:

- If $Fac(n)$ is started with an argument n that is the product of two prime numbers, Fac will output $p_1 > 1, p_2 > 1$ with $p_1 \cdot p_2 = n$.

Safety and liveness

Example: Factorization algorithm $Fac(n) \rightarrow \{p_i\}$:

- If $Fac(n)$ is started with an argument n that is the product of two prime numbers, Fac will output $p_1 > 1, p_2 > 1$ with $p_1 \cdot p_2 = n$.
- Safety property or liveness property?

Safety and liveness

Verification usually is easier if you use multiple minimalistic properties!

Example:

- *Liveness*: The algorithm $Fac(n)$, started with an arbitrary integer n as argument, eventually terminates.
- *Safety*: If $Fac(n)$ terminates and produces output $\{p_1, \dots, p_k\}$, then $\forall i \in \{1..k\} : p_i$ is a prime number > 1 and $\prod_{i=1}^k p_i = n$

Overview

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- Reliable broadcast (rbcast)
- FIFO broadcast (fbcast)
- Causal broadcast (cbcast)
- Atomic broadcast (abcast)

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches

Broadcasts

Fault-Tolerant Broadcasts

- Very fundamental abstraction for distributed systems
- A large variety of broadcasts with different **semantics**
- Algorithms highly depend on **system model**

Broadcasts

Fault-Tolerant Broadcasts

- Very fundamental abstraction for distributed systems
- A large variety of broadcasts with different **semantics**
- Algorithms highly depend on **system model**

This chapter is heavily based

- on the publication “*A modular approach to fault-tolerant broadcasts and related problems*” by Vassos Hadzilacos and Sam Toueg (1994)
 - All following definitions (green boxes) are direct citations from the publication; you should read the article at least up to page 16.
- and on RSDP

Broadcasts: layered architecture

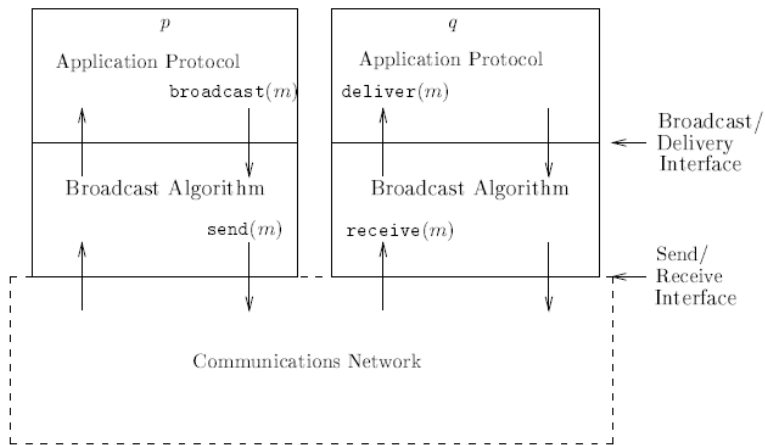


Figure 3: Application/Broadcast Layering

Broadcast semantics

- Best-effort broadcast
- **rbcast**: Reliable broadcast
- **fbcast**: FIFO broadcast
- **cbcast**: Causal broadcast
- **abcast**: Atomic (total-order) broadcast
- **fabcast**: FIFO atomic broadcast
- **cabcast**: Causal atomic broadcast
- Timed broadcast
- Uniform broadcast

Best-effort broadcast

Implementation

- Send message to all nodes using perfect point-to-point links
- (Or use existing low-level broadcast mechanisms, such as Ethernet broadcast)

Characteristics

- Simple, low latency
- No message order guarantees
- Reliability?

Best-effort broadcast

Implementation

- Send message to all nodes using perfect point-to-point links
- (Or use existing low-level broadcast mechanisms, such as Ethernet broadcast)

Characteristics

- Simple, low latency
- No message order guarantees
- Reliable delivery if sender does not crash
 - No guarantees if sender crashes! (→ see next slide)

Best-effort broadcast

Best-effort broadcast – P denotes the set of all processes:

Uses: `PerfectPointToPointLink (ppp)`


```
Upon event < BEBroadcast, m > do:  
    forall p in P do:  
        trigger < ppp.Send, p, m >;
```

```
Upon event < ppp.Receive, m > do:  
    trigger < BEDeliver, m >;
```

Best-effort broadcast

Simple algorithm (best-effort broadcast):


- Works fine as long as sender does not crash
- If sender crashes, some but not all processes might receive message



Best-effort broadcast

Simple algorithm (best-effort broadcast):

- Works fine as long as sender does not crash
- If sender crashes, some but not all processes might receive message
- In some situations, this can be a problem . . .



Overview

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- **Reliable broadcast (rbcast)**
- FIFO broadcast (fbcast)
- Causal broadcast (cbcast)
- Atomic broadcast (abcast)

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches

Reliable broadcast

Basic idea (very informal!):

- Make sure that messages are not lost: all (nonfaulty) processes shall deliver all messages
- No guarantees on message order

Reliable broadcast

Basic idea (very informal!):

- Make sure that messages are not lost: all (nonfaulty) processes shall deliver all messages
- No guarantees on message order
- Before defining a “correct” algorithm, let’s be a bit more formal. . .

Reliable broadcast (rbcast)

- Sender p executes operation $broadcast(m)$ with message m
 - m contains sender ID and unique sequence number of the sender
 - i.e. $m := (sender, seq\#, data)$
- Message m is delivered at q : Operation $deliver(m)$ is executed

Properties of rbcast

- **Validity:** correct $p : broadcast(m) \Rightarrow p : deliver(m)$ within finite time
- **Agreement:** \exists correct $p : deliver(m) \Rightarrow \forall$ correct $q : deliver(m)$ within finite time
- **Integrity:** Each messages m is delivered only once by each correct process, and only if $sender(m)$ did broadcast the message

Reliable broadcast (rbcast)

- Think about an algorithm that solves the problem

Reliable broadcast (rbcast)

- Think about an algorithm that solves the problem

Echo algorithm (from paper):

Algorithm for process p :

To execute broadcast(R, m):

send(m) to p

deliver(R, m) occurs as follows:

upon receive(m) do

if p has not previously executed deliver(R, m)

then

send(m) to all neighbors

deliver(R, m)

Figure 11: Reliable Broadcast for Point-to-Point Networks

Reliable broadcast (rbcast)

In our notation:

Uses: PerfectPointToPointLink (ppp)

upon Init:

```
    delivered := empty set;
```

upon event < RBroadcast, m > do:

```
    trigger < ppp.Send, self, m >;
```

upon event < ppp.Receive, m > do:

```
    if(not delivered.contains(m)) then
```

```
        forall p in P do:
```

```
            trigger < ppp.Send, p, m >;
```

```
        delivered.add(m);
```

```
        trigger < RDeliver, m >;
```


Reliable broadcast (rbcast)

Correctness verification (informal):

- Validity: $\text{RBroadcast} \rightarrow \text{Send} \rightarrow \text{Receive} \rightarrow \text{RDeliver}$ within finite time (assuming that ppp channel delivers message within finite time)
- Agreement: If some correct process trigger RDeliver , it has triggered ppp.Send for all destinations before. Thus, assuming the properties of perfect links, each correct process p_i will trigger the corresponding ppp.Receive
 - If m is not in delivered at p_i , p_i will RDeliver the message
 - If m is in delivered, it was added to delivered before, which implies RDeliver of m
- Integrity:
 - the `delivered` set prevents duplication
 - ppp.Receive happens only after ppp.Send (property of ppp link), which in turn happens only after RBroadcast

Uniform reliable broadcast

- The reliable broadcast so far places no restrictions on messages **delivered by faulty processes**
 - For instance the *agreement* property allows faulty processes to deliver different messages
 - Problematic especially if processes interact with external entities
- Properties can be strengthened (\rightarrow **uniformity**):

Modified properties of uniform rbcst

- **Uniform agreement (einheitliche Übereinstimmung):**
 $\exists p : deliver(m) \Rightarrow \forall \text{ correct } q : deliver(m) \text{ within finite time}$
- **Uniform integrity (einheitliche Integrität):**
Each messages m is delivered only once, and only if m was sent by $sender(m)$ before.

Uniform reliable broadcast

Previous echo algorithm:

- Satisfies uniform integrity
- Does **not** satisfy uniform agreement!

Uniform reliable broadcast

Previous echo algorithm:

- Satisfies uniform integrity
- Does **not** satisfy uniform agreement!

Note (without proof): It is impossible to implement uniform reliable broadcast in an asynchronous system model with an arbitrary number of faulty processes!

- Solutions exist for synchronous model
- Solutions exist for asynchronous model if a majority of processes are correct

Overview

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- Reliable broadcast (rbcast)
- **FIFO broadcast (fbcast)**
- Causal broadcast (cbcast)
- Atomic broadcast (abcast)

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches

FIFO broadcast (fbcast)

Difference: order of messages

- Reliable broadcasts: All correct processes deliver all messages, but order may be completely different (i.e., no order guarantees at all)
- FIFO: If P_i broadcasts m_1 before m_2 , then no correct process will deliver m_2 unless it has previously delivered m_1
- Causal: If m_1 causally precedes m_2 , then no correct process will deliver m_2 unless it has previously delivered m_1
- Total: If two correct processes, A and B , deliver two messages m_1 and m_2 , then B delivers m_2 after m_1 if and only if A delivers m_2 after m_1

FIFO broadcast (fbcast)

Difference: order of messages

- Reliable broadcasts: All correct processes deliver all messages, but order may be completely different (i.e., no order guarantees at all)
- FIFO: If P_i broadcasts m_1 before m_2 , then no correct process will deliver m_2 unless it has previously delivered m_1
- Causal: If m_1 causally precedes m_2 , then no correct process will deliver m_2 unless it has previously delivered m_1
- Total: If two correct processes, A and B , deliver two messages m_1 and m_2 , then B delivers m_2 after m_1 if and only if A delivers m_2 after m_1
- (and combinations FIFO+Total, Causal+Total)

FIFO broadcast (fbcast)

- Add another property to the definition of reliable broadcast:

Additional property of fbcast

- **FIFO order:** If a process broadcasts a message m before it broadcasts a message m' , then no correct process delivers m' unless it has previously delivered m .
- How to implement it?
- We present a *transformation* of reliable broadcast (any algorithm) into FIFO broadcast

FIFO broadcast (fbcast)

Uses: Reliable broadcast (rbcst)

upon event Init:

 msgSet := \emptyset

 next[s] := 1 for each process s

upon event < FBroadcast, m >:

 trigger < rbcst.RBroadcast, m >;

upon event < rbcst.RDeliver, m >:

 s := sender(m);

 if next[s] = seq#(m) then

 trigger < Fdeliver, m >;

 next[s] := next[s] + 1;

 while ($\exists m'$ in msgSet: sender(m')=s and next[s]=seq#(m')):

 trigger < Fdeliver, m' >;

 next[s] := next[s] + 1;

 else

 msgSet := msgSet \cup m;


FIFO broadcast (fbcast)

The same algorithm, less formal:

- Every message has a per-sender sequence number
- If the “right” message arrives: deliver it immediately
- If a message arrives out of order (at least one message with smaller sequence number is missing): Put message in queue and deliver it only after receiving and delivering the missing messages


rbcast vs. fbcast

■ Reliable broadcast



rbcast vs. fbcast

■ FIFO broadcast



■ P3 delivers $m2$ after $m1$

Overview

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- Reliable broadcast (rbcast)
- FIFO broadcast (fbcast)
- **Causal broadcast (cbcast)**
- Atomic broadcast (abcast)

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches

Causal broadcast (cbcast)

Causal broadcast is a reliable broadcast with satisfies the following additional property:

Properties of rbcast

- **Causal order:** If the broadcast of a message m causally precedes the broadcast of a message m' (i.e., if $m \rightarrow m'$), then no correct process delivers m' unless it has previously delivered m .

Causal broadcast (cbcast)

Causal broadcast is a reliable broadcast with satisfies the following additional property:

Properties of rbcast

- **Causal order:** If the broadcast of a message m causally precedes the broadcast of a message m' (i.e., if $m \rightarrow m'$), then no correct process delivers m' unless it has previously delivered m .

- ⇒ See lecture on logical clocks for a definition of causal precedence ($m_1 \rightarrow m_2$)
- ⇒ Note that causal order implies FIFO order

Causal broadcast: simple algorithm

Uses: FIFO broadcast (fbcast)

upon event Init:

 rcntDeliver := \emptyset

upon event $\langle \text{CBroadcast}, m \rangle$:

 trigger $\langle \text{fbcast.FBroadcast}, \langle \text{rcntDeliver} \parallel m \rangle \rangle$;

 rcntDeliver := \emptyset

upon event $\langle \text{fbcast.FDeliver}, \langle m_1, \dots, m_q \rangle \rangle$:

 for $i := 1 \dots q$:

 if p has not previously executed $\text{Cdeliver}(m_i)$:

 trigger $\langle \text{Cdeliver}, m_i \rangle$

 rcntDeliver := rcntDeliver $\parallel m_i$


Causal broadcast: simple algorithm

Basic idea:

- Uses FIFO broadcast \Rightarrow guarantees correct FIFO order
- Records all local deliveries
- Instead of sending m , FIFO-broadcasts all locally delivered messages since last broadcast \Rightarrow guarantees correct order

fbcast vs. cbcast


■ FIFO broadcast



■ P3 delivers $m2$ before $m1$, which violates causal order

fbcast vs. cbcast

■ causal broadcast



■ P3 delivers $m2$ after $m1$

Causal broadcast algorithms

- Previous algorithm involves a considerable overhead: each broadcast includes all causally preceding messages delivered since last broadcast
 - Benefit: algorithm is non-blocking
- Removing this overhead
 - Messages take only the IDs of those in which they depend (largest ID per process)
 - This is equivalent to labelling messages with a *vector clock*
 - Inconvenient: algorithm is blocking

Overview

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- Reliable broadcast (rbcast)
- FIFO broadcast (fbcast)
- Causal broadcast (cbcast)
- **Atomic broadcast (abcast)**

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches

Atomic broadcast (abcasts)

Atomic broadcast or total-order broadcast: defined in terms of

- Reliable broadcast properties +
- Total order property

Additional property of abcast

Total order: If correct processes p and q both deliver messages m and m' , then p delivers m before m' only if q delivers m before m'

Atomic broadcast

■ Algorithm?

Atomic broadcast

- Algorithm?
- There's no (deterministic) algorithm for atomic broadcast in a system model:
 - Asynchronous communication
 - A single (a priori unknown) node may fail (crash fault)
- !? (see later)

More broadcasts

- Atomic FIFO broadcast
 - Combination of atomic broadcast with FIFO order
 - Note: Plain atomic broadcast does *not* necessarily imply FIFO order!
- Atomic causal broadcast
 - Combination of atomic broadcast with causal order

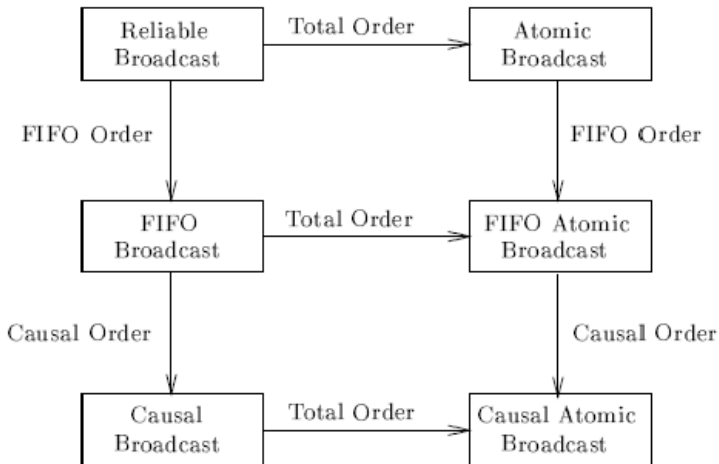
Summary of specifications

- *Validity*: If a correct process broadcasts a message m , then it eventually delivers m .
- *Agreement*: If a correct process delivers a message m , then all correct processes eventually deliver m .
- *Integrity*: For any message m , every correct process delivers m at most once, and only if m was previously broadcast by $sender(m)$.
- *FIFO Order*: If a process broadcasts a message m before it broadcasts a message m' , then no correct process delivers m' unless it has previously delivered m .
- *Causal Order*: If the broadcast of a message m causally precedes the broadcast of a message m' , then no correct process delivers m' unless it has previously delivered m .
- *Total Order*: If correct processes p and q both deliver messages m and m' , then p delivers m before m' if and only if q delivers m before m' .

Summary of specifications

- Reliable broadcast = Validity + Agreement + Integrity
- FIFO broadcast = Reliable broadcast + FIFO order
- Causal broadcast = Reliable broadcast + Causal order
- Atomic broadcast = Reliable broadcast + Total order
- FIFO atomic broadcast = FIFO broadcast + Total order
- Causal atomic broadcast = Causal broadcast + Total order

Summary of specifications



Overview

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- Reliable broadcast (rbcast)
- FIFO broadcast (fbcast)
- Causal broadcast (cbcast)
- Atomic broadcast (abcast)

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches

Reliable broadcast with Byzantine faults

(First) Problem: Faulty sender may send inconsistent values

- Simple echo algorithms (for crashes) makes sure that some value is delivered
- But not that all delivered values are equal
- . . . think about a solution!

Reliable broadcast with Byzantine faults

(First) Problem: Faulty sender may send inconsistent values

- Simple echo algorithms (for crashes) makes sure that some value is delivered
- But not that all delivered values are equal
- . . . think about a solution!

(First) Solution: Collect votes from all processes

- There are at least $n - f$ correct processes
- Accept value v if more than half of all correct processes vote in favour of v
- More than $f + \frac{n-f}{2} = \frac{n+f}{2}$ votes implies votes from more than half of all correct processes (without knowing which processes are correct)

Reliable broadcast with Byzantine faults

Implication (without synchrony):

- f faulty processes might not send anything

Reliable broadcast with Byzantine faults

Implication (without synchrony):

- f faulty processes might not send anything
- Any algorithm should do something after receiving messages from only $n - f$ processes

Reliable broadcast with Byzantine faults

Implication (without synchrony):

- f faulty processes might not send anything
- Any algorithm should do something after receiving messages from only $n - f$ processes
- If sender is correct, a positive vote should always be made, i.e., $n - f > \frac{n+f}{2}$ must hold
- This is equivalent to $n > 3f$

Reliable broadcast with Byzantine faults

Implication (without synchrony):

- f faulty processes might not send anything
- Any algorithm should do something after receiving messages from only $n - f$ processes
- If sender is correct, a positive vote should always be made, i.e., $n - f > \frac{n+f}{2}$ must hold
- This is equivalent to $n > 3f$

For tolerating f Byzantine faults, we (usually) need at least $3f + 1$ processes

Echo broadcast with Byzantine faults

- $n = 3f + 1$ processes (or generically $f = \lfloor (n - 1)/3 \rfloor$)
- What does it guarantee?

- 1 G sends (initial, G, m) message to all processes.
- 2 When a process receives the first (initial, G, m) message from G, it sends an (echo, G, m) message to all processes. All subsequent (initial, G, m') messages are ignored.
- 3 If a process receives an (echo, G, m) message from more than $\frac{n+f}{2}$ distinct processes, then it accepts the message m from G.

(from: Sam Toueg: Randomized Byzantine Agreements, 1984)

Echo broadcast with Byzantine faults

(in Toueg-Paper: correct = “proper”, deliver = “accept”)

If a process G broadcasts the message m with echo broadcast then:

- ① The messages delivered by correct processes are identical
- ② If G is correct, then all the correct processes deliver m

- *Note: Property 1 guarantees only that the delivered messages are identical (i.e., if two messages are delivered, they are the same). If G is faulty, it is possible that some processes deliver m , and others do **not** deliver m .*

Reliable broadcast with Byzantine faults

For reliable broadcast:

- Requirement: If one correct delivers some message m , all correct processes must eventually deliver m (even if $sender(m)$ is faulty)

Reliable broadcast with Byzantine faults

For reliable broadcast:

- Requirement: If one correct delivers some message m , all correct processes must eventually deliver m (even if $sender(m)$ is faulty)
- Algorithm by Bracha 1984
 - Use “more than $\frac{n+f}{2}$ votes” argument to ensure that at most one message can be delivered.
 - Use additional step to guarantee “all-or-nothing” behaviour if sender is faulty.

Reliable broadcast (Bracha 1984)

■ $n = 3f + 1$, operation Broadcast(v)

① (by transmitter only): send (initial, v) to all processes.

② Wait until receive for some v :

- one (initial, v) message
- or $> \frac{n+f}{2}$ (echo, v) messages
- or $(f + 1)$ (ready, v) messages

send (echo, v) to all processes

⑤ Wait until receive for some v :

- $> \frac{n+f}{2}$ (echo, v) messages
- or $(f + 1)$ (ready, v) messages

send (ready, v) to all proc.

⑥ Wait until receive for some v :

- $2f + 1$ (ready, v) messages

accept v

(from: G. Bracha: An asynchronous $\lfloor (n - 1)/3 \rfloor$ -resilient consensus protocol, 1984)

Reliable broadcast (Bracha 1984): Properties

- 1 If p is correct, then all the correct processes accept the value of its message.
- 2 If p is malicious, then either all the correct processes accept the same value, or non of them will accept any value from p .

Reliable broadcast (Bracha 1984): Proof

- ❶ **Lemma 1:** If two correct processes r and s send $(ready, v)$ and $(ready, u)$ messages, respectively, then $u = v$.
- ❷ **Lemma 2:** If two correct processes p and q accept the values v and u , respectively, then $u == v$.
- ❸ **Lemma 3:** If a correct process p accepts the value v , then every other correct process will eventually accept v .
- ❹ **Lemma 4:** If the transmitter p is correct and it sends v , then all correct processes will accept v .

Summary of broadcast echo algorithms

Crash faults

- Simple echo

Byzantine faults

- Simple echo + threshold
- “Two-phase echo” + thresholds

Summary of broadcast echo algorithms

Crash faults

- Simple echo for implementing **reliable broadcast**

Byzantine faults

- Simple echo + threshold for implementing **best effort broadcast**
- “Two-phase echo” + thresholds for implementing **reliable broadcast**

Overview

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- Reliable broadcast (rbcast)
- FIFO broadcast (fbcast)
- Causal broadcast (cbcast)
- Atomic broadcast (abcast)

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches

Implementing atomic broadcast

Classification of well-known algorithms according to Défago et al.:

- Sequencer-based approaches
 - Selected node determines message order
 - Fixed or moving sequencer
- Token-based approaches (“privilege-based”)
 - Serialisation of all messages: only token owner may send
 - Active token / passive token
- “Communication history”
 - Sender assigns time stamps, receiver sorts messages
- Consensus-based approaches
 - Order determined by distributed consensus algorithms

Sequencer-based approaches

Sequencer defines message order

■ Variant “UB”:

- Sender sends message to sequencer (unicast)
- Sequencer sends message to all nodes (broadcast)
- High load / bottleneck at sequencer

■ Variant “UUB”:

- Sender sends message to sequencer (unicast)
- Sequencer assigns sequence number to sender (unicast)
- Sender sends message + sequence number to all nodes (broadcast)
- Less work for sequencer, but higher latency

■ Variant “BB”:

- Sender sends message to all nodes (including sequencer; broadcast)
- Sequencer assigns sequence number and sends it to all nodes (broadcast)

In all cases: What if sequencer crashes?

Additionally for UUB and BB: What if sender crashes?

Token-based approaches

- Only token owner may send messages
- Circulating token
 - Unused token is passed to next node automatically
 - Efficient if many nodes broadcast messages
 - Token circulates even if no one wants to broadcast messages
- Requesting the token on demand
 - Token is passed only upon request
 - Efficient if only one / few nodes want to broadcast
 - Many senders: High overhead for searching for the token

Communication history-based approaches

- Order is defined by the sender
- Messages usually carry a (physical or logical) timestamp
- Recipients observe messages generated by other processes and their timestamp
 - i.e., the “communication history”
- Process can deliver a message m as soon as it knows that there cannot be any other message that needs to be delivered before m

Communication history-based approaches

- Order is defined by the sender
- Messages usually carry a (physical or logical) timestamp
- Recipients observe messages generated by other processes and their timestamp
 - i.e., the “communication history”
- Process can deliver a message m as soon as it knows that there cannot be any other message that needs to be delivered before m

Simple examples:

- Synchronous system: Physical time stamps; delivery delayed by upper bound of clock error + transmission delay
- Asynchronous system: Round-robin approach: Deliver message 1 of P_1, P_2, \dots , then message 2, etc.

Consensus-based approaches

Distributed consensus algorithms:

- Each node proposes some value
In this case: next message ID / list of n message IDs
- Result of consensus: One of the proposed values Here: Next message / next n ordered messages

Next: transformation: abcast \leftrightarrow consensus

Atomic broadcast and consensus

So far

- Simple solutions for bcast, fbcast, cbcast in asynchronous system
- Solution for uniform bcast if less than half of processes is faulty
- No solution for abcast in asynchronous system

Next step:

- Atomic broadcast and the consensus problem

Consensus

Definition of the *consensus problem*:

Each process (in a set of n processes) proposes a value, and all correct processes decide upon a common value, such that the following properties are satisfied:

- *Termination*: Every correct process eventually decides some value
- *Agreement*: No two correct processes decide on different values v, v'
- *Validity*: If a correct process decides v , then v was previously proposed by some process.
- *Integrity*: No process decides twice

Note: some other definitions with minor differences exist!

Consensus variants

- *Uniform agreement*: If a process (whether correct or faulty) decides v , then all correct processes eventually decide v .
- *Uniform integrity*: If a process (whether correct or faulty) decides v , then v was previously proposed by some process.
- *Validity'*: If all processes propose the same value v , then this is the only possible decision value
- *Validity''*: If all correct processes propose the same value v , then this is the only possible decision value

Transforming atomic broadcast into consensus

- Algorithm for atomic broadcast also solves the consensus problem. . .
 - The transformation uses no timing assumption, so it works for any synchrony model

Uses: Atomic Broadcast (abcast)

upon Init:

 decided := false;

upon event < Propose, v > do:

 trigger < abcast.Broadcast, v >;

upon event < abcast.Deliver, u > do:

 if not decided then

 trigger < Decide, u >;

Transforming rbcast + consensus into abcast

Uses: Reliable Broadcast (rbcast), MultiConsensus (cons)

upon Init:

 R_delivered := \emptyset ; A_delivered := \emptyset ; k := 0

upon event < ABroadcast, m > do:

 trigger < rbcast.RBroadcast, m >;

upon event < rbcast.RDeliver, m > do:

 R_delivered := R_delivered \cup {m},

repeat forever:

 A_undelivered := R_delivered - A_delivered

 if A_undelivered \neq 0 then

 k := k + 1

 trigger < cons_k.Propose, A_undelivered >

 wait

upon event < cons_k.Decide, msgSet >:

 forall v in (msgSet - A_delivered) in deterministic order

 trigger < ADeliver, v >

 A_delivered := A_delivered \cup msgSet

 resume wait

Summary

1 Basics: Safety and Liveness

2 Broadcast

- Best-effort broadcast
- Reliable broadcast (rbcast)
- FIFO broadcast (fbcast)
- Causal broadcast (cbcast)
- Atomic broadcast (abcast)

3 Reliable broadcast with Byzantine faults

4 Implementing atomic broadcast

- Sequencer-based approaches
- Token-based approaches
- Communication history-based approaches
- Consensus-based approaches