


Dependable Distributed Systems – 5880V/UE

Part 8. Distributed File Systems – 2021-10-14

Prof. Dr. Hans P. Reiser | WS 2021/22

UNIVERSITÄT PASSAU



(source: Ghemawat et al.: The Google File System, SOSP 2013)

Today's class

What concepts allow you to store large amount of data in distributed systems?

Today's class

What concepts allow you to store large amount of data in distributed systems?

Requirements

- Reliability
- Security
- Consistency
- . . .

Today's class

What concepts allow you to store large amount of data in distributed systems?

Requirements

- Reliability
- Security
- Consistency
- ...

Context

- in closely coupled, large-scale data centers

Overview

- 1 Large-scale file systems
 - Example: Google File System (GFS)

High volume distributed file systems

Examples:

- Google File System – GFS
 - Slides mostly based on the GFS SOSP'03 presentation
- HDFS (Yahoo!/Apache)
 - Similar, but things have different names, some simplifications
 - Open source
 - Used by: Amazon/A9, Facebook, Google, Joost, Last.fm, New York Times, PowerSet, Veoh, Yahoo!, . . .

Introduction

Design constraints:

- Component failures are the norm
 - 1000s of components
 - Bugs, human errors, failures of memory, disk, connectors, networking, and power supplies
 - Mechanisms: monitoring, error detection, fault masking, automatic recovery
- Files are huge by traditional standards
 - Multi-GB files are common
 - *Important design option: modest number of huge files*

Introduction

Design constraints:

- Most modifications are appends
 - Random writes are practically nonexistent
 - Many files are written once, and read sequentially
- Two types of reads
 - Large streaming reads
 - Small random reads
- Sustained throughput more important than latency
- File system APIs are open to changes

Interface Design

- Not POSIX compliant
- Standard operations
 - create, delete, open, close, read, write
- Additional operations
 - Snapshot (cheap copy of a file or a directory)
 - Record append (allows concurrent atomic appends from multiple clients)

Architectural Design

A GFS cluster:

- A single **master** → contains *metadata* only
- Multiple **chunkservers** per master → contain *data*
 - Accessed by multiple clients
- Running on commodity Linux machines

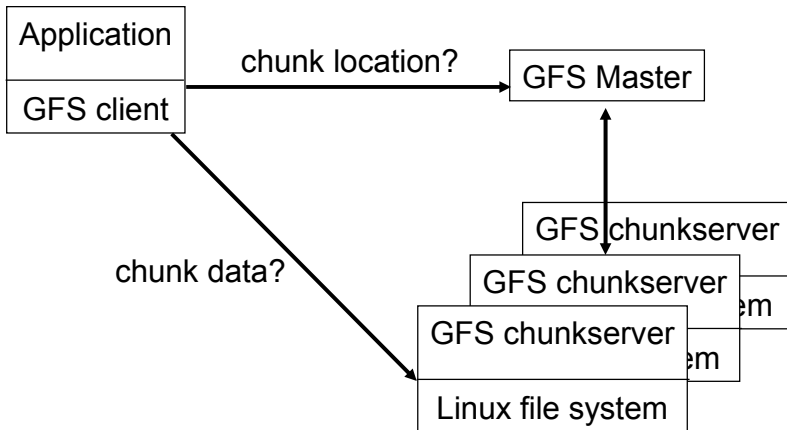
A file:

- Represented as fixed-sized **chunks**
 - Labeled with 64-bit unique global IDs
 - Stored at chunkservers
 - Replicated in several chunkservers (3 by default)

File x

chunk 1
chunk 2
chunk 3
...
chunk n

Architectural Design (2)



Architectural Design (3)

Master server:

- Maintains all metadata
 - Name space, access control, file-to-chunk mappings, garbage collection, chunk migration
- Bottleneck, so they try to minimize its workload / functionality

GPS clients:

- Consult master for metadata
- Access data from chunkservers
- No caching at clients and chunkservers due to the frequent case of streaming (well, almost)

Single-Master Design

- Simple
- Master answers only chunk locations
- A client typically asks for multiple chunk locations in a single request
- The master also predictively provides chunk locations immediately following those requested

Chunk Size

- 64 MB
- Fewer chunk location requests to the master
- Reduced overhead to access a chunk
- Persistent TCP connection to the chunkserver over an extended period of time
- Fewer metadata entries
 - Kept in memory
- Potential problem: Hot spot development

Metadata

Three major types

- a) File and chunk namespaces — persistent
- b) File-to-chunk mappings — persistent
- c) Locations of a chunk's replicas — non persistent

Metadata

All kept in memory

- Fast!
- Quick global scans
 - Garbage collections
 - Reorganizations
- < 64 bytes per 64 MB of data
- Prefix compression
- But some data is also logged

Master fault tolerance (1)

Locations of a chunk's replicas is not persistent — (c)

- Master polls chunkservers at startup
 - “What chunks do you have?”
- Use heartbeat messages to monitor chunkservers
- Simplicity
- On-demand approach vs. coordination
 - On-demand wins when changes (failures) are frequent

Master fault tolerance (2)

- File and chunk namespaces and File-to-chunk mappings — (a),(b)
- Variation of **passive replication**
- Metadata updates are **logged**
 - Log replicated on remote machines
- Take global **snapshots (checkpoints)** to truncate logs
 - Memory mapped (no serialization/deserialization)
 - Checkpoints can be created while updates arrive
- Recovery
 - Latest checkpoint + subsequent log files

Consistency Model

Namespace mutations (e.g., file creations / deletions) are atomic

- Total order, handled exclusively by the master

Relaxed consistency

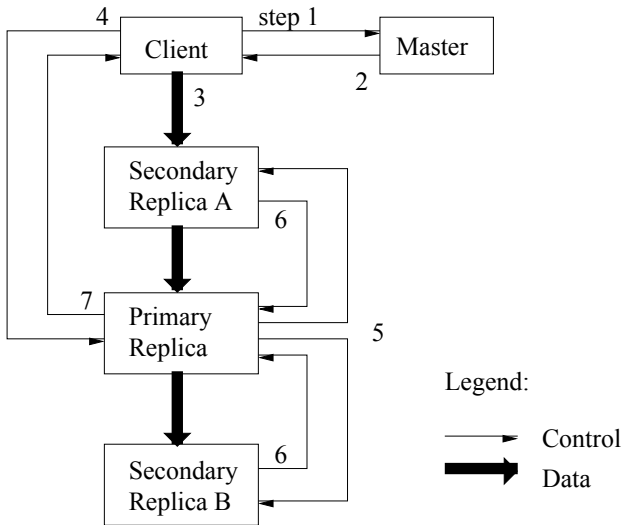
- Concurrent changes to a chunk are:
 - Consistent — all clients see the same data
 - Undefined — the chunk may end with a mix of the writes done
- An append is atomically committed at least once
 - But GFS defines the offset (i.e., position of the append at the file)
- All changes to a chunk are applied in the same order to all replicas
- Use version numbers to detect missed updates

System Interactions

(objective: spare the master)

- The master grants a **chunk lease** to a replica
 - Replica = chunkserver that has the chunk
- The replica holding the lease determines the order of updates to all replicas
 - The replica becomes the **primary** for that chunk
- Lease
 - 60 second timeouts
 - Can be extended indefinitely
 - Extension request piggybacked on heartbeat messages
 - After timeout expires, master can grant new leases

Write control and data flow



Data Flow

- Separation of control and data flows
 - Avoid network bottleneck
- Updates are pushed linearly among replicas
- Pipelined transfers
- Instead of sending data around as needed, GFS “thinks” about the paths it should follow in order to minimize bottlenecks

Snapshot of a file

- Copy-on-write approach
- When the Master receives a file snapshot request:
 - (1) Revoke all leases on the file chunks
 - ⇒ all subsequent updates to the file must pass through the Master
 - ⇒ Master can do a copy of the file first (any updates are logged while taking the snapshot)
 - (2) Replicate the metadata but leave it pointing to the same chunks
 - (3) Apply the log to a copy of the metadata
 - A chunk is not copied until the next update (copy-on-write)

Master Operation

Master:

- Executes all namespace operations
- Manages chunk replicas throughout the system
 - Makes placement decisions
 - Creates new chunks and replicas
 - Coordinates activities to keep chunks fully replicated, to balance load across all the chunkservers, and to reclaim unused storage

Some characteristics:

- Does not have a data structure per-directory
- No hard links and symbolic links
- Full path name to metadata mapping
 - With prefix compression

Locking Operations

Locking a file

- To lock /d1/d2/leaf
- Need to lock /d1, /d1/d2, and /d1/d2/leaf
- Totally ordered locking to prevent deadlocks

Replica Placement

- *Replicas of chunks*
- Goals:
 - Maximize data reliability and availability
 - Maximize network bandwidth
- Need to spread chunk replicas across machines and racks
- Higher priority to replica chunks with lower replication factors
- Limited resources spent on replication

Garbage Collection of files

- Garbage collection is simpler than eager deletion due to
 - Unfinished replica creation
 - Lost deletion messages
- Deleted files are hidden for three days
- Then they are garbage collected
- Combined with other background operations (taking snapshots)
- Safety net against accidents

Fault Tolerance and Diagnosis

- Fast recovery
 - Master and chunkserver are designed to restore their states and start in seconds regardless of how they terminated (normal, abnormal, killed, . . .)
- Chunk replication
- Master replication
 - *Shadow masters* provide read-only access when the primary master is down
 - If master fails it is restarted immediately
 - If it cannot be restarted (due to hardware failure), a new one is elected from outside GFS (Chubby?)

Fault Tolerance and Diagnosis

Data integrity

- A chunk is divided into 64-KB blocks
- Each with its checksum
- Verified at read and write times
- Also background scans for rarely used data
 - Locally at each chunkserver

Measurements

- Chunkserver workload
 - Bimodal distribution of small and large files
 - Ratio of write to append operations: 3:1 to 8:1
 - Virtually no overwrites
- Master workload
 - Most request for chunk locations and open files
- Reads achieve 75% of the network limit
- Writes achieve 50% of the network limit

Major Innovations

- File system API tailored to specific workload
- Single-master design to simplify coordination
- Metadata fit in memory
- Flat namespace